
Formatting Instructions For NeurIPS 2023

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 The abstract paragraph should be indented 1/2 inch (3 picas) on both the left- and
2 right-hand margins. Use 10 point type, with a vertical spacing (leading) of 11 points.
3 The word **Abstract** must be centered, bold, and in point size 12. Two line spaces
4 precede the abstract. The abstract must be limited to one paragraph.

5 1 Submission of papers to NeurIPS 2023

6 2 Introduction

7 3 Diffusion Models

8 In this section we quickly review the basics of diffusion models. We focus on the stochastic differential
9 equation formulation first presented by (author?) [5].

10 Let $p(\mathbf{y})_{\text{data}}$ denote the data distribution. The goal of a diffusion model is to learn a mapping from a
11 simple distribution $p(\mathbf{z})$ to the data distribution $p(\mathbf{y})_{\text{data}}$.

12 This is achieved by reversing a diffusion process. In particular, we construct a stochastic differential
13 equation $\mathbf{y}(t)$ from $t \in [0, T]$ such that $\mathbf{y}(0) \sim p(\mathbf{y})_{\text{data}}$ and $\mathbf{y}(1) \sim p(\mathbf{y}(T))$ is a simple distribution
14 we can sample from and whose evolution is given by

$$d\mathbf{y}(t) = \mathbf{f}(\mathbf{y}(t), t)dt + \mathbf{g}(t)d\mathbf{w}(t), \quad (1)$$

15 where $\mathbf{w}(t)$ is a standard Brownian motion and $\mathbf{f} : \mathbb{R}^d \times [0, T] \rightarrow \mathbb{R}^d$ and $\mathbf{g} : [0, T] \rightarrow \mathbb{R}$ are called
16 the drift coefficient and the diffusion coefficient, respectively.

17 It is possible to reverse this SDE and sample from $p(\mathbf{y})_{\text{data}}$ by first sampling $\mathbf{y}(1) \sim p(\mathbf{y}(T))$ and
18 then evolving the system backwards in time. This is done by solving the reverse SDE (Cite anderson
19 1982)

$$d\mathbf{y}(t) = [f(\mathbf{y}(t), t) - g(t)^2 \nabla_{\mathbf{y}(t)} \log p(\mathbf{y}(t))]dt + g(t)d\bar{\mathbf{w}}(t). \quad (2)$$

20 where $\bar{\mathbf{w}}(t)$ is a standard Brownian with reversed time. Thus because \mathbf{f} and g are known, and we
21 construct the SDE so that $p(\mathbf{y}(T))$ is simple, as long as we know the score $\nabla_{\mathbf{y}(t)} \log p(\mathbf{y}(t))$ we can
22 sample from $p(\mathbf{y})_{\text{data}}$.

23 Estimating the Score

24 An important result by (author?) [6] is that it is possible to estimate the score $\nabla_{\mathbf{y}(t)} \log p(\mathbf{y}(t))$ by
25 computing

$$\mathbf{s}^* = \operatorname{argmin}_{\mathbf{s} \in \mathcal{S}} \mathbb{E}_t \mathbb{E}_{p(\mathbf{y}(0))_{\text{data}}} \mathbb{E}_{p(\mathbf{y}(t)|\mathbf{y}(0))} \left[\left\| \nabla_{\mathbf{y}(t)} \log p(\mathbf{y}(t)|\mathbf{y}(0)) - \mathbf{s}(\mathbf{y}(t), t) \right\|^2 \right]. \quad (3)$$

where $\mathcal{S} = \{s : \mathbb{R}^d \times [0, T] \rightarrow \mathbb{R}^d\}$ is the set of all possible score functions indexed by time t , and \mathbb{E}_t denotes the expectation over uniformly sampled $t \in [0, T]$.

Conditional Diffusion Models

Although it is most common to train diffusion models unconditionally as explained above, one can also train diffusion models conditionally on some input \mathbf{x} .

To do so we make the following modifications to the above formulation.

1. We construct one separate SDE per value of \mathbf{x} . Each SDE shares the same drift and diffusion coefficients but the initial distribution $p_{\mathbf{x}}(\mathbf{y}(0))$ is given by $p(\mathbf{y}|\mathbf{x})_{\text{data}}$.
2. The reverse SDE is now given by

$$d\mathbf{y}(t) = [f(\mathbf{y}(t), t) - g(t)^2 \nabla_{\mathbf{y}(t)} \log p_{\mathbf{x}}(\mathbf{y}(t))]dt + g(t)d\bar{\mathbf{w}}(t). \quad (4)$$

where $\nabla_{\mathbf{y}(t)} \log p_{\mathbf{x}}(\mathbf{y}(t))$ is the score of the conditional distribution $p_{\mathbf{x}}(\mathbf{y}(t))$. Importantly, because we choose the diffusion and drift coefficients so that at $t = T$ the distribution is the same for all values of \mathbf{x} , we can still sample from the data distribution in the same way as before.

3. The final change is that the score function is now estimated by

$$\mathbf{s}^* = \operatorname{argmin}_{\mathbf{s} \in \mathcal{S}} \mathbb{E}_t \mathbb{E}_{p(\mathbf{y}(0), \mathbf{x})_{\text{data}}} \mathbb{E}_{p(\mathbf{y}(t)|\mathbf{y}(0))} \left[\|\nabla_{\mathbf{y}(t)} \log p(\mathbf{y}(t)|\mathbf{y}(0)) - \mathbf{s}(\mathbf{y}(t), \mathbf{x}, t)\|^2 \right].$$

with the changes being that now $\mathcal{S} = \{s : \mathbb{R}^d \times \mathbb{R}^m \times [0, T] \rightarrow \mathbb{R}^d\}$ is the set of all possible score functions but now allowing for the score to depend on the input \mathbf{x} , and the expectation is taken over the joint distribution $p(\mathbf{y}(0), \mathbf{x})_{\text{data}}$. We emphasize that the score of the conditional distribution $p(\mathbf{y}(t)|\mathbf{y}(0))$ is still the same because once we condition on $\mathbf{y}(0)$ the distribution is the same for all values of \mathbf{x} .

This formulation of conditional diffusion models is different than controllable generation as presented in [5]. There, a conditional diffusion model is constructed by noting that

$$\nabla_{\mathbf{y}(t)} p(\mathbf{y}(t)|\mathbf{x}) = \nabla_{\mathbf{y}(t)} p(\mathbf{y}(t)) + \nabla_{\mathbf{y}(t)} p(\mathbf{x}|\mathbf{y}(t))$$

and hence if we obtain the first term from an unconditional diffusion model, and the second term by differentiating through another trained model $p(\mathbf{x}|\mathbf{y}(t))$, we can obtain the score of the conditional distribution. In our case this is not feasible because in general the dimension of \mathbf{y} will be much smaller than the dimension of \mathbf{x} .

4 Gradient Boosted Trees

Gradient Boosted Trees (GBT) [2] are a popular non-parametric machine learning model for function approximation. The objective is to find a function $F : \mathbb{R}^d \rightarrow \mathbb{R}$ that minimizes

$$L(F) = \mathbb{E}_{\mathbf{x}, y} [l(y, F(\mathbf{x}))], \quad (5)$$

where $l : \mathbb{R} \times \mathbb{R} \rightarrow \mathbb{R}$ is a loss function, and the expectation is taken over the joint distribution of the input \mathbf{x} and the target y . It does this by imposing the requirement that F is as a scaled sum of M decision trees $f_m : \mathbb{R}^d \rightarrow \mathbb{R}$, i.e.

$$F(\mathbf{x}) = \sum_{m=1}^M \epsilon f_m(\mathbf{x}), \quad \epsilon \in (0, 1). \quad (6)$$

where ϵ is a learning rate or shrinkage parameter. In the most basic form of the algorithm each tree is constructed to approximate gradient descent on the loss function $L(F)$. In particular, if we let $F_i = \sum_{m=1}^i f_m$ denote the function after i iterations and then the i -th tree is constructed to approximately minimize the squared error

$$f_i = \operatorname{argmin}_f \mathbb{E}_{\mathbf{x}, y} \left(f(\mathbf{x}) - \frac{\partial l(y, \hat{y})}{\partial \hat{y}} \Big|_{\hat{y}=F_i(\mathbf{x})} \right)^2. \quad (7)$$

61 using empirical risk minimization and a greedy algorithm to construct the tree.
 62 Various modifications to the basic algorithm have been proposed and implemented such as reg-
 63 ularization, special ways of optimizing the tree, support for categorical functions, higher order
 64 optimization[1, 3, 4]. In this paper we focus on the implementation of GBTs in the LightGBM library
 65 [3] with the understanding that the same principles apply to any other GBT implementation.

66 5 Treeffuser/Treeffusion Models

67 For $\mathbf{x} \in \mathbb{R}^d$, $\mathbf{y} \in \mathbb{R}^m$ the objective of probabilistic predictions is to produce an estimate of the full
 68 conditional distribution $\mathbb{P}[\mathbf{y}|\mathbf{x}]$. This objective is different than in standard regression where the goal
 69 is usually to predict $\mathbb{E}[\mathbf{y}|\mathbf{x}]$.

70 The most common approach to solve this problem is via parametric models. This procedure assumes
 71 that the distribution $\mathbb{P}[\mathbf{y}|\mathbf{x}]$ can be well approximated by a parametric family of distributions

$$\mathbb{P}[\mathbf{y}|\mathbf{x}] = p[\mathbf{y}|\theta(\mathbf{x})],$$

72 where p is a well known distribution (e.g Gaussian) and $\theta(\mathbf{x})$ is a function that maps \mathbf{x} to the
 73 parameters of the distribution p (e.g. the mean and covariance of a Gaussian). Optimization is then
 74 performed by finding the function $\theta(\mathbf{x})$ that minimizes a proper-scoring rule such as the negative
 75 log-likelihood.

76 References

77 References follow the acknowledgments in the camera-ready paper. Use unnumbered first-level
 78 heading for the references. Any choice of citation style is acceptable as long as you are consistent. It
 79 is permissible to reduce the font size to small (9 point) when listing the references. Note that the
 80 Reference section does not count towards the page limit.(author?) [5]

81 References

- 82 [1] Tianqi Chen and Carlos Guestrin. Xgboost: A scalable tree boosting system. In *Proceedings of*
 83 *the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*,
 84 KDD '16. ACM, August 2016.
- 85 [2] Jerome H Friedman. Greedy function approximation: a gradient boosting machine. *Annals of*
 86 *Statistics*, 29(5):1189–1232, 2001.
- 87 [3] Guolin Ke, Qi Meng, Thomas Finley, Taifeng Wang, Wei Chen, Weidong Ma, Qiwei Ye, and
 88 Tie-Yan Liu. Lightgbm: A highly efficient gradient boosting decision tree. In I. Guyon, U. Von
 89 Luxburg, S. Bengio, H. Wallach, R. Fergus, S. Vishwanathan, and R. Garnett, editors, *Advances*
 90 *in Neural Information Processing Systems*, volume 30. Curran Associates, Inc., 2017.
- 91 [4] Liudmila Prokhorenkova, Gleb Gusev, Aleksandr Vorobev, Anna Veronika Dorogush, and Andrey
 92 Gulin. Catboost: unbiased boosting with categorical features, 2019.
- 93 [5] Yang Song, Jascha Sohl-Dickstein, Diederik P. Kingma, Abhishek Kumar, Stefano Ermon, and
 94 Ben Poole. Score-based generative modeling through stochastic differential equations, 2021.
- 95 [6] Pascal Vincent. A connection between score matching and denoising autoencoders. Technical
 96 Report 1358, Département d’Informatique et de Recherche Opérationnelle, Université de Mon-
 97 tréal, CP 6128, Succ. Centre-Ville, Montréal (QC) H3C 3J7, Canada, December 2010. THIS IS
 98 A PREPRINT VERSION OF A NOTE THAT HAS BEEN ACCEPTED FOR PUBLICATION
 99 IN NEURAL COMPUTATION.